

```
# Libraries to help with reading and manipulating data
import numpy as np
import pandas as pd

# Libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

sns.set()

# Removes the limit for the number of displayed columns
pd.set_option("display.max_columns", None)
# Sets the limit for the number of displayed rows
pd.set_option("display.max_rows", 200)

# to split the data into train and test
from sklearn.model_selection import train_test_split

# to build linear regression_model
from sklearn.linear_model import LinearRegression

# to check model performance
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# I changed this part
!pip install mlxtend
import joblib
import sys
sys.modules['sklearn.externals.joblib'] = joblib
from mlxtend.feature_selection import SequentialFeatureSelector as SFS
```

```
Requirement already satisfied: mlxtend in /usr/local/lib/python3.10/dist-packages (0.22.0)
Requirement already satisfied: scipy>=1.2.1 in /usr/local/lib/python3.10/dist-packages (from mlxtend) (1.11.4)
Requirement already satisfied: numpy>=1.16.2 in /usr/local/lib/python3.10/dist-packages (from mlxtend) (1.25.2)
Requirement already satisfied: pandas>=0.24.2 in /usr/local/lib/python3.10/dist-packages (from mlxtend) (2.0.3)
Requirement already satisfied: scikit-learn>=1.0.2 in /usr/local/lib/python3.10/dist-packages (from mlxtend) (1.2.2)
Requirement already satisfied: matplotlib>=3.0.0 in /usr/local/lib/python3.10/dist-packages (from mlxtend) (3.7.1)
Requirement already satisfied: joblib>=0.13.2 in /usr/local/lib/python3.10/dist-packages (from mlxtend) (1.4.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from mlxtend) (67.7.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0.0->mlxtend) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24.2->mlxtend) (2023.4)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24.2->mlxtend) (2024.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.0.2->mlxtend) (3.4.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib>=3.0.0->mlxtend) (1.16.0)
```

```
# Importing Libraries
import requests
import pandas as pd
from imblearn.over_sampling import SMOTE
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn import svm
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB as Naive_Bayes
from sklearn import metrics
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.datasets import make_classification
from xgboost import XGBClassifier
from sklearn.externals import joblib

from IPython.display import display
import pickle
```

```
df = pd.read_csv('/content/parkinsons dataset.csv')
```

Double-click (or enter) to edit

```
df.head()
```

	name	MDVP:F0(Hz)	MDVP:F1(Hz)	MDVP:F2(Hz)	MDVP:Jitter(%)	MDVP:Jitter(Abs)	MDVP:RAP	MDVP:PPQ	Jitter:DDP	MDVP:Shimmer
0	phon_R01_S01_1	119.992	157.302	74.997	0.00784	0.00007	0.00370	0.00554	0.01109	0.043
1	phon_R01_S01_2	122.400	148.650	113.819	0.00968	0.00008	0.00465	0.00696	0.01394	0.061
2	phon_R01_S01_3	116.682	131.111	111.555	0.01050	0.00009	0.00544	0.00781	0.01633	0.052
3	phon_R01_S01_4	116.676	137.871	111.366	0.00997	0.00009	0.00502	0.00698	0.01505	0.054
4	phon_R01_S01_5	116.014	141.781	110.655	0.01284	0.00011	0.00655	0.00908	0.01966	0.064

```
df.tail()
```

	name	MDVP:F0(Hz)	MDVP:F1(Hz)	MDVP:F2(Hz)	MDVP:Jitter(%)	MDVP:Jitter
191	phon_R01_S50_3	209.516	253.017	89.488	0.00564	
192	phon_R01_S50_4	174.688	240.005	74.287	0.01360	
193	phon_R01_S50_5	198.764	396.961	74.904	0.00740	
194	phon_R01_S50_6	214.289	260.277	77.973	0.00567	
195	phon_R01_S50_6	216.289	266.279	77.974	0.00568	

```
print('Number of Features In Dataset :', df.shape[1])
print('Number of Instances In Dataset : ', df.shape[0])
```

```
Number of Features In Dataset : 24
Number of Instances In Dataset : 196
```

```
# Dropping The Name Column
df.drop(['name'], axis=1, inplace=True)
```

```
print('Number of Features In Dataset :', df.shape[1])
print('Number of Instances In Dataset : ', df.shape[0])
```

Number of Features In Dataset : 23
 Number of Instances In Dataset : 196

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 196 entries, 0 to 195
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   MDVP:Fo(Hz)           196 non-null    float64
1   MDVP:Fhi(Hz)          196 non-null    float64
2   MDVP:Flo(Hz)          196 non-null    float64
3   MDVP:Jitter(%)        196 non-null    float64
4   MDVP:Jitter(Abs)      196 non-null    float64
5   MDVP:RAP               196 non-null    float64
6   MDVP:PPQ              196 non-null    float64
7   Jitter:DDP            196 non-null    float64
8   MDVP:Shimmer          196 non-null    float64
9   MDVP:Shimmer(dB)      196 non-null    float64
10  Shimmer:APQ3          196 non-null    float64
11  Shimmer:APQ5          196 non-null    float64
12  MDVP:APQ              196 non-null    float64
13  Shimmer:DDA           196 non-null    float64
14  NHR                   196 non-null    float64
15  HNR                   196 non-null    float64
16  status                196 non-null    int64
17  RPDE                  196 non-null    float64
18  DFA                   196 non-null    float64
19  spread1               196 non-null    float64
20  spread2               196 non-null    float64
21  D2                    196 non-null    float64
22  PPE                   196 non-null    float64
dtypes: float64(22), int64(1)
memory usage: 35.3 KB
```

```
df.describe()
```

	MDVP:Fo(Hz)	MDVP:Fhi(Hz)	MDVP:Flo(Hz)	MDVP:Jitter(%)	MDVP:Jitter(Abs)	MDVP
count	196.000000	196.000000	196.000000	196.000000	196.000000	196.00
mean	154.545276	197.457847	116.128964	0.006218	0.000044	0.00:
std	41.521110	91.390318	43.496022	0.004836	0.000035	0.00:
min	88.333000	102.145000	65.476000	0.001680	0.000007	0.00:
25%	117.721000	134.965750	84.044250	0.003460	0.000020	0.00
50%	149.239500	176.212000	104.205000	0.004945	0.000030	0.00:
75%	183.653750	224.804250	139.504250	0.007347	0.000060	0.00:
max	260.105000	592.030000	239.170000	0.033160	0.000260	0.02

```
df['status'] = df['status'].astype('uint8')
```

```
# Checking For Duplicate Rows In Dataset
print('Number of Duplicated Rows :',df.duplicated().sum())
```

Number of Duplicated Rows : 0

```
# Checking For Missing Values In Dataset
df.isna().sum()
```

```
MDVP:Fo(Hz)      0
MDVP:Fhi(Hz)     0
MDVP:Flo(Hz)     0
MDVP:Jitter(%)   0
MDVP:Jitter(Abs) 0
MDVP:RAP         0
MDVP:PPQ         0
Jitter:DDP       0
MDVP:Shimmer     0
MDVP:Shimmer(dB) 0
Shimmer:APQ3     0
```

```

Shimmer:APQ5      0
MDVP:APQ          0
Shimmer:DDA       0
NHR               0
HNR               0
status            0
RPDE              0
DFA               0
spread1           0
spread2           0
D2                0
PPE               0
dtype: int64


```

```

#Balance of Data
sns.countplot(x='status',data=df)

```

<Axes: xlabel='status', ylabel='count'>

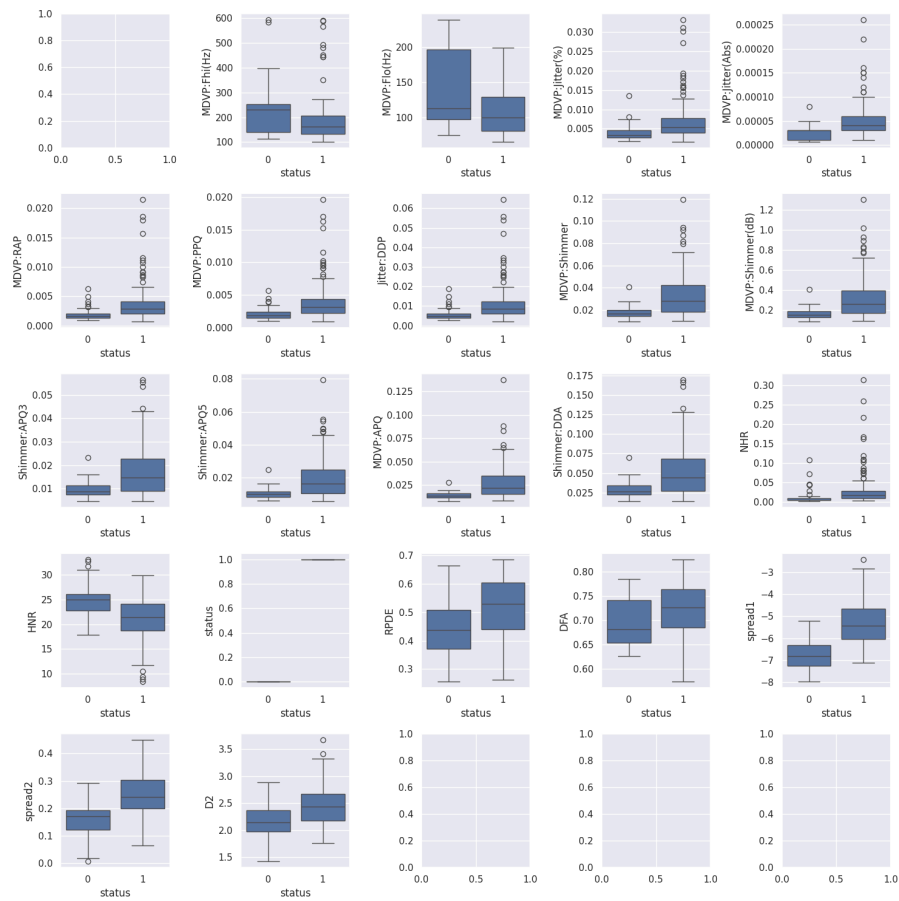


```


#Box Plot
fig,axes=plt.subplots(5,5,figsize=(15,15))
axes=axes.flatten()

for i in range(1,len(df.columns)-1):
    sns.boxplot(x='status',y=df.iloc[:,i],data=df,orient='v',ax=axes[i])
plt.tight_layout()
plt.show()

```




```
plt.rcParams['figure.figsize'] = (15, 4)
sns.pairplot(df, hue = 'status', vars = ['MDVP:Jitter(%)', 'MDVP:Jitter(Abs)', 'MDVP:RAP', 'MDVP:PPQ', 'Jitter:DDP'] )
plt.show()
```

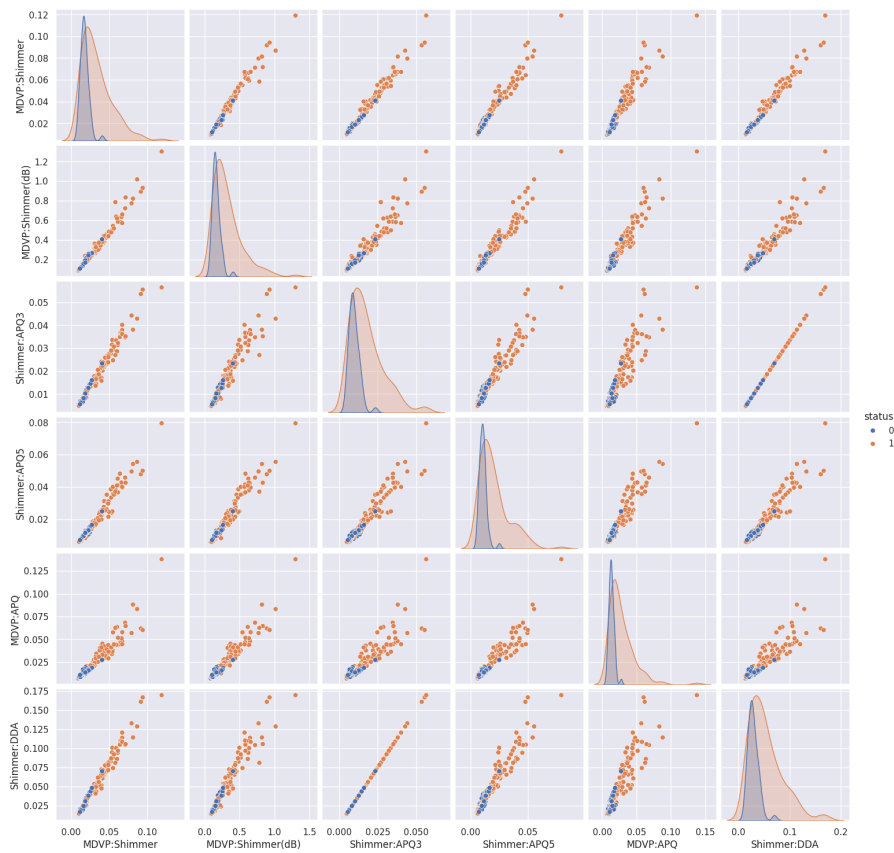


```
fig, ax = plt.subplots(figsize=(20,20))
sns.heatmap(df.corr(), annot=True, ax=ax)
```

<Axes: >



```
plt.rcParams['figure.figsize'] = (15, 4)
sns.pairplot(df,hue = 'status', vars = ['MDVP:Shimmer','MDVP:Shimmer(dB)','Shimmer:APQ3','Shimmer:APQ5','MDVP:APQ','Shimmer:DDA'] )
plt.show()
```



```
# Exploring Imbalance In Dataset
```

```
df['status'].value_counts()
```

```
status
1    147
0     49
Name: count, dtype: int64
```

```
X = df.drop(['status'], axis=1)
```

```
y = df['status']
```

```
print('Feature (X) Shape Before Balancing :', X.shape)
```

```
print('Target (y) Shape Before Balancing :', y.shape)
```

```
Feature (X) Shape Before Balancing : (196, 22)
```

```
Target (y) Shape Before Balancing : (196,)
```

```
y=df['status']
```

```
cols=['MDVP:RAP', 'Jitter:DDP', 'DFA', 'NHR', 'MDVP:Fhi(Hz)', 'status']
```

```
x=df.drop(cols,axis=1)
```



```
from sklearn.preprocessing import MinMaxScaler
```

```
# Scaling features between -1 and 1 for normalization
scaler = MinMaxScaler((-1,1))
```

```
# define X_features , Y_labels
X_features = scaler.fit_transform(X)
Y_labels = y
```

```
# splitting the dataset into training and testing sets into 80 - 20
train_size=0.80
test_size=0.20
seed=5
from sklearn.model_selection import train_test_split
X_train , X_test , y_train , y_test = train_test_split(X_features, Y_labels , test_size=0.20, random_state=20)
```


```
logmodel = LogisticRegression()
logmodel.fit(X_train, y_train)
predlog = logmodel.predict(X_test)
```

```
print(classification_report(y_test, predlog))
print("Confusion Matrix:")
confusion_matrix(y_test, predlog)
```

	precision	recall	f1-score	support
0	0.86	0.46	0.60	13
1	0.79	0.96	0.87	27
accuracy			0.80	40
macro avg	0.82	0.71	0.73	40
weighted avg	0.81	0.80	0.78	40

```
Confusion Matrix:
array([[ 6,  7],
       [ 1, 26]])
```

```
y_pred_proba = logmodel.predict_proba(X_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



```
# Dumping Logistic Regression Model
pickle.dump(logmodel, open('lg.pkl', 'wb'))
```

```
# Naive Bayes
```

```
gnb = Naive_Bayes()
gnb.fit(X_train, y_train)
predgnb = gnb.predict(X_test)

print(classification_report(y_test, predgnb))
```

	precision	recall	f1-score	support
0	0.50	1.00	0.67	13
1	1.00	0.52	0.68	27
accuracy			0.68	40
macro avg	0.75	0.76	0.67	40
weighted avg	0.84	0.68	0.68	40


```
print("Confusion Matrix:")
confusion_matrix(y_test, predgnb)
```

```
Confusion Matrix:
array([[13,  0],
       [13, 14]])
```

```
# scores -check how efficiently labels are predicted
accuracy_testing = accuracy_score(y_test, predgnb)
print("Accuracy % :",accuracy_testing*100)
```

```
Accuracy % : 67.5
```

```
y_pred_proba = gnb.predict_proba(X_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



```
pickle.dump(gnb,open ('gnb.pkl','wb'))
```

```
import numpy as np
```

```
Ks = 10
```

```
mean_acc = []
```

```
ConfusionMx = [];
```

```
for n in range(2,Ks):
```

```
    #Train Model and Predict
```

```
    neigh = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
```

```
    yhat=neigh.predict(X_test)
```

```
    mean_acc.append(metrics.accuracy_score(y_test, yhat))
```

```
print('Neighbor Accuracy List')
```

```
print(mean_acc)
```

```
Neighbor Accuracy List
```

```
[0.925, 0.95, 0.925, 0.85, 0.85, 0.85, 0.825, 0.825]
```

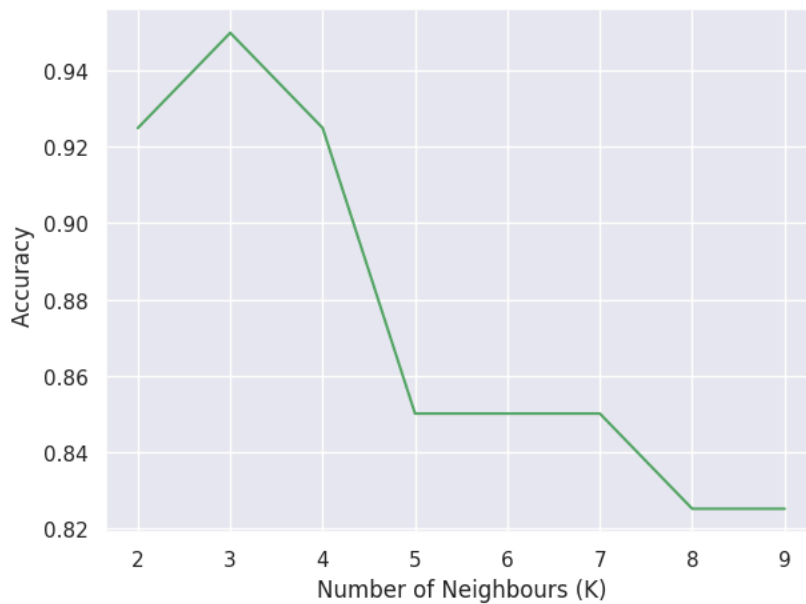
```
plt.plot(range(2,Ks),mean_acc,'g')
```

```
plt.ylabel('Accuracy ')
```

```
plt.xlabel('Number of Neighbours (K)')
```

```
plt.tight_layout()
```

```
plt.show()
```



```
knn = KNeighborsClassifier(n_neighbors=5)
```

```
knn.fit(X_train, y_train)
```

```
predKNN = knn.predict(X_test)
```

```
y_pred_proba = knn.predict_proba(X_test)[::,1]
```


```
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
```

```
auc = metrics.roc_auc_score(y_test, y_pred_proba)
```

```
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
```

```
plt.legend(loc=4)
```

```
plt.show()
```



```
# Dumping KNN Classifier
pickle.dump(knn, open('knn_clf.pkl', 'wb'))
```

```
# Defining Parameter Dictionary
param_dict = {'max_depth': range(4,8), 'eta' : [0.1, 0.2, 0.3, 0.4, 0.5],
              'reg_lambda' : [0.8, 0.9, 1, 1.1, 1.2],
              'random_state': [300, 600, 900]}
```

```
from sklearn.metrics import precision_score, recall_score, accuracy_score, f1_score, r2_score, log_loss
```

```
chart = {
    'Metric': ["Accuracy", "F1-Score", "Recall", "Precision", "R2-Score"],
    'LR': [accuracy_score(y_test, predlog), f1_score(y_test, predlog), recall_score(y_test, predlog), precision_score(y_test, predlog), r2_score(y_test, predlog)],
    'NB': [accuracy_score(y_test, predgnb), f1_score(y_test, predgnb), recall_score(y_test, predgnb), precision_score(y_test, predgnb), r2_score(y_test, predgnb)],
    'KNN': [accuracy_score(y_test, predKNN), f1_score(y_test, predKNN), recall_score(y_test, predKNN), precision_score(y_test, predKNN), r2_score(y_test, predKNN)],
}
chart = pd.DataFrame(chart)
```


```
display(chart)
```

	Metric	LR	NB	KNN
0	Accuracy	0.800000	0.675000	0.850000
1	F1-Score	0.866667	0.682927	0.892857
2	Recall	0.962963	0.518519	0.925926
3	Precision	0.787879	1.000000	0.862069
4	R2-Score	0.088319	-0.481481	0.316239

▼ NB

```
# @title NB
```

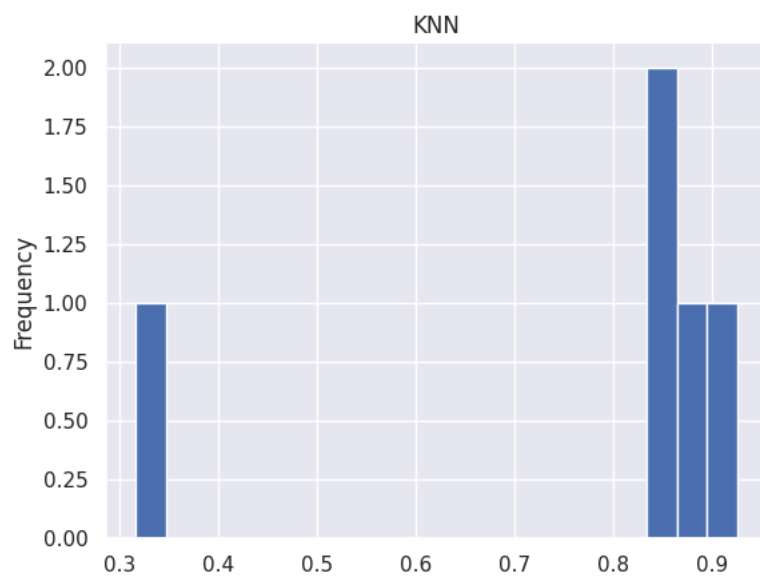
```
from matplotlib import pyplot as plt
chart['NB'].plot(kind='hist', bins=20, title='NB')
plt.gca().spines[['top', 'right']].set_visible(False)
```



▼ KNN

@title KNN


```
from matplotlib import pyplot as plt
chart['KNN'].plot(kind='hist', bins=20, title='KNN')
plt.gca().spines[['top', 'right']].set_visible(False)
```



▼ LR

@title LR

```
from matplotlib import pyplot as plt
chart['LR'].plot(kind='hist', bins=20, title='LR')
plt.gca().spines[['top', 'right']].set_visible(False)
```



```

from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Initialize Logistic Regression and Naive Bayes classifiers
logmodel = LogisticRegression()
gnb = GaussianNB()

# Fit both models on training data
logmodel.fit(X_train, y_train)
gnb.fit(X_train, y_train)

# Predict probabilities for both models
log_prob = logmodel.predict_proba(X_test)[: , 1]
gnb_prob = gnb.predict_proba(X_test)[: , 1]

# Weighted average of probabilities
hybrid_prob = (log_prob + gnb_prob) / 2

# Compute accuracy for each model
log_accuracy = accuracy_score(y_test, logmodel.predict(X_test))
gnb_accuracy = accuracy_score(y_test, gnb.predict(X_test))

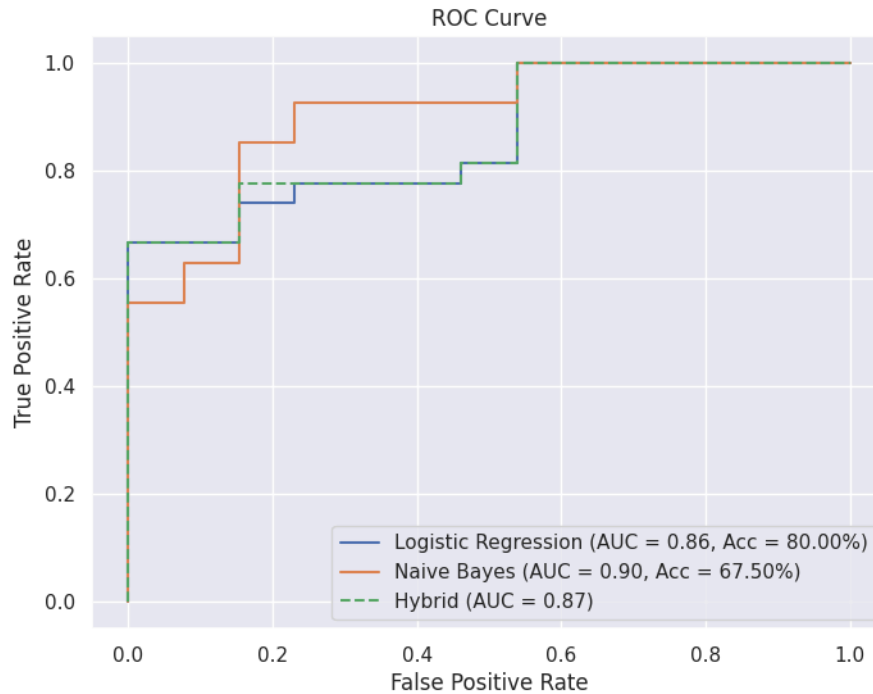
# Compute AUC for each model
log_auc = roc_auc_score(y_test, log_prob)
gnb_auc = roc_auc_score(y_test, gnb_prob)
hybrid_auc = roc_auc_score(y_test, hybrid_prob)

# Plot ROC curves for each model
plt.figure(figsize=(8, 6))
fpr_log, tpr_log, _ = roc_curve(y_test, log_prob)
fpr_gnb, tpr_gnb, _ = roc_curve(y_test, gnb_prob)
fpr_hybrid, tpr_hybrid, _ = roc_curve(y_test, hybrid_prob)

plt.plot(fpr_log, tpr_log, label='Logistic Regression (AUC = {:.2f}, Acc = {:.2f}%)'
        .format(log_auc, log_accuracy * 100))
plt.plot(fpr_gnb, tpr_gnb, label='Naive Bayes (AUC = {:.2f}, Acc = {:.2f}%)'
        .format(gnb_auc, gnb_accuracy * 100))
plt.plot(fpr_hybrid, tpr_hybrid, label='Hybrid (AUC = {:.2f})'
        .format(hybrid_auc), linestyle='--')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

```



```

from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Initialize Naive Bayes and KNN classifiers
gnb = GaussianNB()
knn = KNeighborsClassifier(n_neighbors=5)

# Fit both models on training data
gnb.fit(X_train, y_train)
knn.fit(X_train, y_train)

# Predict probabilities for both models
gnb_prob = gnb.predict_proba(X_test)[: , 1]
knn_prob = knn.predict_proba(X_test)[: , 1]

# Weighted average of probabilities
hybrid_prob = (gnb_prob + knn_prob) / 2

# Compute accuracy for each model
gnb_accuracy = accuracy_score(y_test, gnb.predict(X_test))
knn_accuracy = accuracy_score(y_test, knn.predict(X_test))

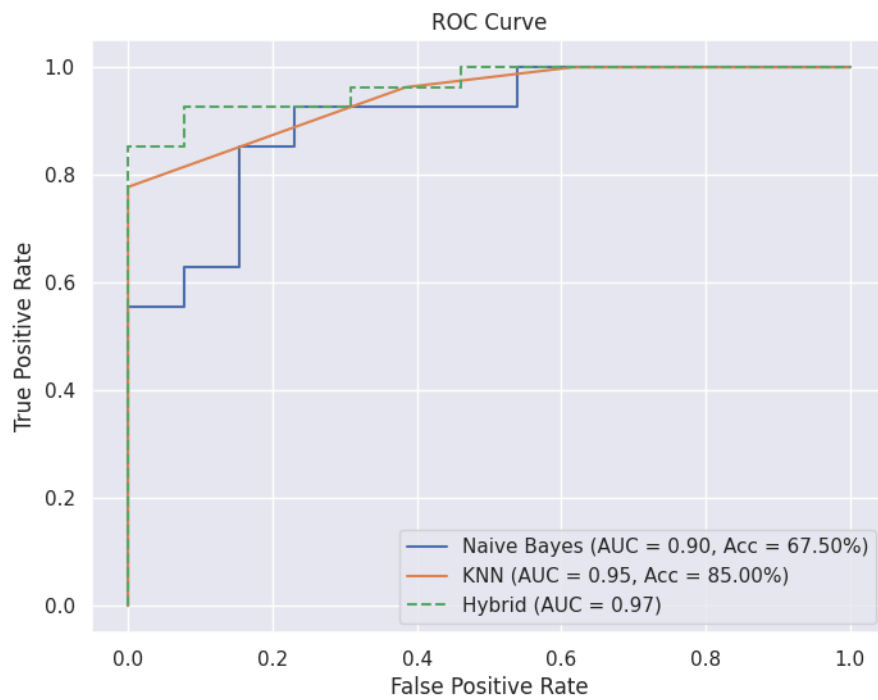
# Compute AUC for each model
gnb_auc = roc_auc_score(y_test, gnb_prob)
knn_auc = roc_auc_score(y_test, knn_prob)
hybrid_auc = roc_auc_score(y_test, hybrid_prob)

# Plot ROC curves for each model
plt.figure(figsize=(8, 6))
fpr_gnb, tpr_gnb, _ = roc_curve(y_test, gnb_prob)
fpr_knn, tpr_knn, _ = roc_curve(y_test, knn_prob)
fpr_hybrid, tpr_hybrid, _ = roc_curve(y_test, hybrid_prob)

plt.plot(fpr_gnb, tpr_gnb, label='Naive Bayes (AUC = {:.2f}, Acc = {:.2f}%)'
        .format(gnb_auc, gnb_accuracy * 100))
plt.plot(fpr_knn, tpr_knn, label='KNN (AUC = {:.2f}, Acc = {:.2f}%)'
        .format(knn_auc, knn_accuracy * 100))
plt.plot(fpr_hybrid, tpr_hybrid, label='Hybrid (AUC = {:.2f})'
        .format(hybrid_auc), linestyle='--')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

```




```

from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score
import matplotlib.pyplot as plt

# Initialize classifiers
logmodel = LogisticRegression()
gnb = GaussianNB()
knn = KNeighborsClassifier(n_neighbors=5)

# Fit all models on training data
logmodel.fit(X_train, y_train)
gnb.fit(X_train, y_train)
knn.fit(X_train, y_train)

# Predict probabilities for all models
log_prob = logmodel.predict_proba(X_test)[:, 1]
gnb_prob = gnb.predict_proba(X_test)[:, 1]
knn_prob = knn.predict_proba(X_test)[:, 1]

# Weighted average of probabilities
hybrid_prob = (log_prob + gnb_prob + knn_prob) / 3

# Compute accuracy for each model
log_accuracy = accuracy_score(y_test, logmodel.predict(X_test))
gnb_accuracy = accuracy_score(y_test, gnb.predict(X_test))
knn_accuracy = accuracy_score(y_test, knn.predict(X_test))


# Compute AUC for each model
log_auc = roc_auc_score(y_test, log_prob)
gnb_auc = roc_auc_score(y_test, gnb_prob)
knn_auc = roc_auc_score(y_test, knn_prob)
hybrid_auc = roc_auc_score(y_test, hybrid_prob)

# Plot ROC curves for each model
plt.figure(figsize=(8, 6))
fpr_log, tpr_log, _ = roc_curve(y_test, log_prob)
fpr_gnb, tpr_gnb, _ = roc_curve(y_test, gnb_prob)
fpr_knn, tpr_knn, _ = roc_curve(y_test, knn_prob)
fpr_hybrid, tpr_hybrid, _ = roc_curve(y_test, hybrid_prob)

plt.plot(fpr_log, tpr_log, label='Logistic Regression (AUC = {:.2f}, Acc = {:.2f}%)'
        .format(log_auc, log_accuracy * 100))
plt.plot(fpr_gnb, tpr_gnb, label='Naive Bayes (AUC = {:.2f}, Acc = {:.2f}%)'
        .format(gnb_auc, gnb_accuracy * 100))
plt.plot(fpr_knn, tpr_knn, label='KNN (AUC = {:.2f}, Acc = {:.2f}%)'
        .format(knn_auc, knn_accuracy * 100))
plt.plot(fpr_hybrid, tpr_hybrid, label='Hybrid (AUC = {:.2f})'
        .format(hybrid_auc), linestyle='--')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

```



```
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming you have a DataFrame 'data' with multiple observations for each algorithm
# For demonstration purposes, I'll create a random dataset
import pandas as pd
import numpy as np

algorithms = ['Logistic Regression', 'Naive Bayes', 'KNN', 'Hybrid']
auc_scores = np.random.rand(100, 4) # Random AUC scores for 100 observations and 4 algorithms

# Create a DataFrame for visualization
data = pd.DataFrame(auc_scores, columns=algorithms)

# Plotting box plots for AUC scores
plt.figure(figsize=(10, 6))
sns.boxplot(data=data)
plt.title('AUC Scores Comparison')
plt.xlabel('Algorithm')
plt.ylabel('AUC')
plt.show()
```